

GNN Toolkit (Python)

code overview and usage

August 9, 2025

Contents

| | |
|---|----------|
| 1 Overview | 1 |
| 2 Quick Start | 2 |
| 3 Configuration (config.py) | 2 |
| 3.1 <code>GNNConfig</code> | 2 |
| 4 Data & Snapshots (data_loading.py) | 3 |
| 4.1 <code>make_snapshots</code> | 3 |
| 4.2 <code>split_snapshots</code> | 3 |
| 4.3 <code>build_dataloaders</code> | 4 |
| 5 Preprocessing (data/preprocess.py) | 4 |
| 5.1 <code>scale_targets</code> | 4 |
| 6 Model Registry & Zoo (models/registry.py, models.py) | 4 |
| 6.1 Registry | 4 |
| 6.2 Built-ins | 4 |
| 7 Losses, Class Weights, Metrics | 4 |
| 7.1 Loss registry (loss/loss_functions.py) | 4 |
| 7.2 Class imbalance (loss/class_weights.py) | 4 |
| 7.3 Metrics (loss/metrics.py) | 5 |
| 8 Training Loop (trainer.py) | 5 |
| 8.1 <code>Trainer</code> | 5 |
| 9 Design Notes | 5 |
| 10 API Index | 5 |

1 Overview

This package builds graph snapshots from Pandas DataFrames and trains node-level models with PyTorch Geometric. It provides:

- a central run config (`GNNConfig`) with YAML/JSON I/O,
- utilities to convert DataFrames \rightarrow `Data` snapshots \rightarrow `DataLoaders`,
- an extensible model zoo (GCN, GraphSAGE, GAT, GATv2, Simple baseline),
- loss & metric registries (BCE/focal BCE/MSE/MAE/Huber, AUC/ACC/MAE/R²),

- optional class imbalance handling and per-node weighting,
- a high-level `Trainer` with early stopping, schedulers, TensorBoard and checkpointing.

Requirements Python 3.10+, torch, torch_geometric, pandas, numpy, scikit-learn (for AUC), optional pyyaml.

2 Quick Start

Listing 1: Load frames, build loaders, train, evaluate

```

1 from jp_da_imb.gnn.config import load_config
2 from jp_da_imb.gnn.data.preprocess import scale_targets
3 from jp_da_imb.gnn.data_loading import build_data loaders
4 from jp_da_imb.gnn.trainer import Trainer
5 import networkx as nx
6
7 # 1) Load settings
8 cfg = load_config("configs/gnn_run.yaml")
9
10 # 2) Prepare node/edge frames dicts keyed by region and "src->dst"
11 node_frames = {"tokyo": df_tokyo, "kansai": df_kansai, ...} # indexed by timestamp
12 edge_frames = {"tokyo->kansai": df_tk, "kansai->tokyo": df_kt, ...} # indexed by timestamp
13
14 # 3) Optional: scale target(s) using training split stats; keep inverse() for reporting
15 scaled_nodes, scaler = scale_targets(node_frames, cfg)
16
17 # 4) Graph topology
18 g = nx.DiGraph()
19 g.add_nodes_from(sorted(node_frames))
20 g.add_edges_from([k.split("->") for k in edge_frames])
21
22 # 5) Dataloaders
23 train_dl, val_dl, test_dl = build_data loaders(
24     node_frames=scaled_nodes, edge_frames=edge_frames, graph=g, cfg=cfg
25 )
26
27 # 6) Train
28 trainer = Trainer(cfg, train_dl, val_dl, test_dl, log_dir=None, ckpt_dir="./ckpts")
29 trainer.fit()
30
31 # 7) Evaluate (+ optional tidy DataFrame)
32 stats, df = trainer.evaluate(split="test", return_df=True)
33 df_inv = scaler.inverse(df) # add 'pred_orig' / 'target_orig' cols (de-normalised)
34 print(stats)

```

3 Configuration (config.py)

3.1 `GNNConfig`

- **Task/model:** task="node_clf", model_name in {gcnn, graphsage, gat, gatv2, simple}, num_layers, hidden_dim, heads, dropout, norm in {batch, layer, None}.
- **Data:** target_col, split_mode in {date, ratio}, cutoff_date, val_ratio, test_ratio, shuffle_in_split.
- **Optimisation:** batch_size, lr, epochs, patience, grad_clip, optimiser in {adam, adamw, sgd} and optimiser_kwargs.

- **Scheduler:** `lr_scheduler` in `{step, plateau, cosine, onecycle}` plus `lr_scheduler_kwargs`; `print_lr_each_epoch`.
- **Loss/metric:** `loss_fn` in `{bce, focal_bce, mse, mae, huber}`; `class_weights` ("auto" / scalar / list); `node_pos_weights` (per node); `metric` in `{acc, auc, mae, r2}`.
- **Misc:** `in_dropout`, `edge_dropout`, `device`, `run_name`, `seed`.

I/O helpers

- `GNNConfig.load(x)` accepts an existing instance, a dict, or a YAML/JSON path.
- `to_dict()`, `to_yaml()`.
- Free function `load_config(x)` mirrors `GNNConfig.load`.

Example YAML

```
task: node_clf
model_name: gatv2
num_layers: 2
hidden_dim: 64
heads: 4
dropout: 0.5
norm: batch
split_mode: date
cutoff_date: 2023-07-01
val_ratio: 0.10
test_ratio: 0.10
batch_size: 32
lr: 1e-3
epochs: 100
patience: 10
optimiser: adamw
lr_scheduler: plateau
lr_scheduler_kwargs: { factor: 0.5, patience: 3 }
loss_fn: focal_bce
class_weights: auto
metric: auc
device: cuda
run_name: demo_run
```

4 Data & Snapshots (data_loading.py)

4.1 `make_snapshots`

Converts aligned DataFrames into an ordered list of `torch_geometric.data.Data`:

- Node order is sorted(`graph.nodes`); edge order follows the graph's edges.
- Each snapshot has: `x` [`N`, `F_node`], `edge_index` [`2`, `E`], `edge_attr` [`E`, `F_edge`] (optional), `y` [`N`, `T`], optional `node_weight` [`N`] or [`N`, `T`], and `snap_time` (int nanoseconds).
- Multi-target: set `cfg.target_cols` (else falls back to `cfg.target_col`).
- Optional per-node weights via `cfg.node_pos_weights`.

4.2 `split_snapshots`

- `split_mode="date"`: uses `cutoff_date` to split; post-cutoff half is val and half is test.

- `split_mode="ratio"`: uses `val_ratio`, `test_ratio`.

4.3 `build_dataloaders`

High-level wrapper: frames \rightarrow snapshots \rightarrow splits \rightarrow 3 `DataLoader`s. Batch size and `shuffle_in_split` come from `GNNConfig`.

5 Preprocessing (`data/preprocess.py`)

5.1 `scale_targets`

Z-scores `target_col` per region using *training* subset only. Returns:

- scaled node frames (same keys);
- a `TargetScaler` that holds per-region μ/σ and `inverse(df)` to add de-scaled columns (`pred_orig`, `target_orig`, or indexed for multi-T).

6 Model Registry & Zoo (`models/registry.py`, `models.py`)

6.1 Registry

`register("name")` decorates a builder; `build_model(name=..., **kw)` instantiates it.

6.2 Built-ins

- `gcn`: wrapper around `torch_geometric.nn.GCN` with `act="relu"`, `norm {batch,layer,None}`.
- `graphsage` (`sage` alias): adapter that calls GraphSAGE with `Data`.
- `gat`: wrapper for GAT.
- `gatv2`: flexible stack of GATv2Conv layers (keeps concatenated heads), optional `edge_attr`, per-layer norm & ReLU, final linear head.
- `simple`/`baseline`: two GCNConv layers + linear head.

7 Losses, Class Weights, Metrics

7.1 Loss registry (`loss/loss_functions.py`)

`build_loss(name=..., class_weights=..., node_reduction=..., **loss_kw)` returns a callable:

- `"bce"`: `binary_cross_entropy_with_logits`, optional global `pos_weight`.
- `"focal_bce"`: focal term $(1 - p_t)^\gamma$ with optional `alpha`.
- `"mse"`, `"mae"`, `"huber"`.
- Per-node weights are broadcast and applied before `"mean"/"sum"` reduction.

7.2 Class imbalance (`loss/class_weights.py`)

`compute_class_weights(train_loader)` makes one pass over training targets and returns

$$\text{pos_weight} = \frac{\# \text{neg} + \varepsilon}{\# \text{pos} + \varepsilon}$$

as a scalar (single target) or vector (per target).

7.3 Metrics (`loss/metrics.py`)

Registry with:

- "auc": ROC AUC on sigmoid(logits).
- "acc": threshold 0.5 on sigmoid(logits).
- "mae": L1 between predictions and targets.
- "r2": $1 - SS_{\text{res}}/SS_{\text{tot}}$ (regression).

8 Training Loop (`trainer.py`)

8.1 Trainer

Initialises device, infers `d_in`, `d_out`, optional `edge_dim` from a training batch, builds the model via the registry, constructs the loss callable (supports `class_weights="auto"`), selects optimiser (adam/adamw/sgd) and LR scheduler (step/plateau/cosine/onecycle).

Features

- `_train_epoch`: standard loop with optional gradient clipping.
- `_eval`: averages loss & metric over a loader.
- `fit`: main loop with scheduler stepping, TensorBoard logging, console prints, early stopping (patience) and `best.pt` checkpoint writing.
- `evaluate(split, return_df=False)`: computes mean loss/metric and optionally returns a tidy DataFrame with rows (`snap_time`, `node`, `pred[#]`, `target[#]`).
- `predict(...)`: (post-fit) forward-only helper analogous to `evaluate`.

9 Design Notes

- Timestamps across all frames are intersected to avoid leakage/misalignment.
- Node order is fixed (`sorted(graph.nodes)`) so tensors align across snapshots.
- Losses are callables to make weighting strategies pluggable without touching training code.
- Config is the single source of truth; YAML/JSON keeps runs reproducible.

10 API Index

| Module | Functions / Classes |
|---------------------|---|
| config | <code>GNNConfig</code> , <code>load_config</code> |
| data_loading | <code>make_snapshots</code> , <code>split_snapshots</code> , <code>build_data loaders</code> |
| data/preprocess | <code>scale_targets</code> , <code>TargetScaler.inverse</code> |
| models/registry | <code>register</code> , <code>build_model</code> |
| models | <code>build_gcn</code> , <code>build_graphsage</code> , <code>build_gat</code> , <code>build_gatv2</code> , <code>build_simple</code> |
| loss/loss_functions | <code>register</code> , <code>build_loss</code> (bce, focal_bce, mse, mae, huber) |
| loss/class_weights | <code>compute_class_weights</code> |
| loss/metrics | <code>register</code> , METRICS (auc, acc, mae, r2) |
| trainer | <code>Trainer.fit</code> , <code>Trainer.evaluate</code> , <code>Trainer.predict</code> |